

Experiment: Working with K3d and Local Persistent Volumes

Note: Refer to the K3D Getting Started Experiment if you haven't already installed k3d for installation.

Make the local directory that we'll mount into our k3d containers.

For Windows:

Open a command prompt in "Run as administrator" mode



Command Prompt

App



Open



Run as administrator

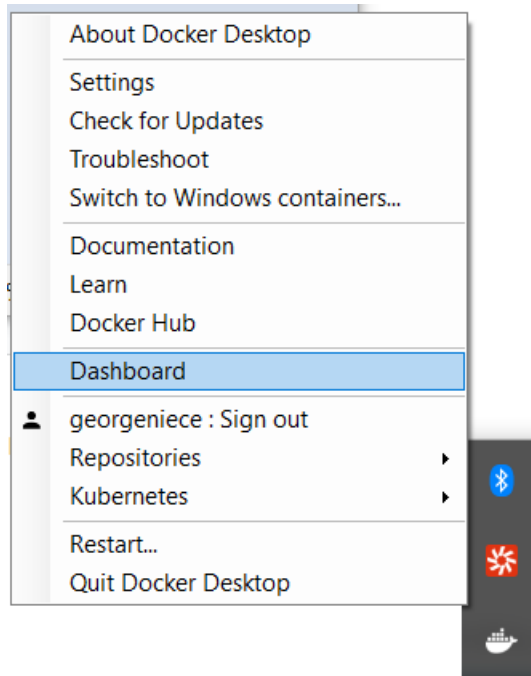
Make the folders that we'll mount into our containers for this experiment

mkdir c:\tmp

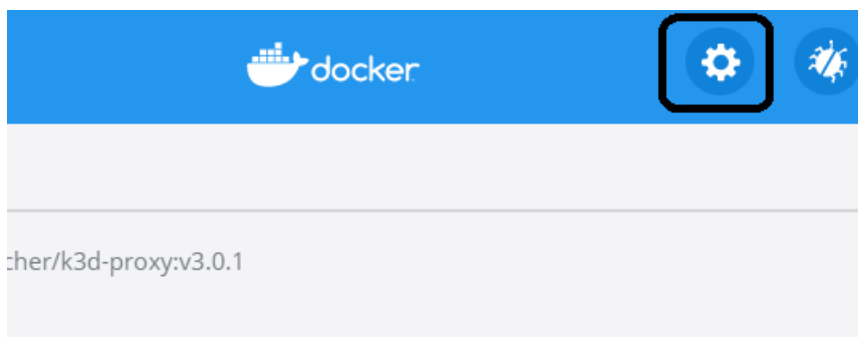
mkdir c:\tmp\k3dvol

The following step is only required if you're using a version of K3D prior to the WSL2 update. You can verify by opening the Docker for Desktop. If there is a "File Sharing" under resources then you would still need to do the following, but it is recommended to update your Docker Desktop before proceeding.

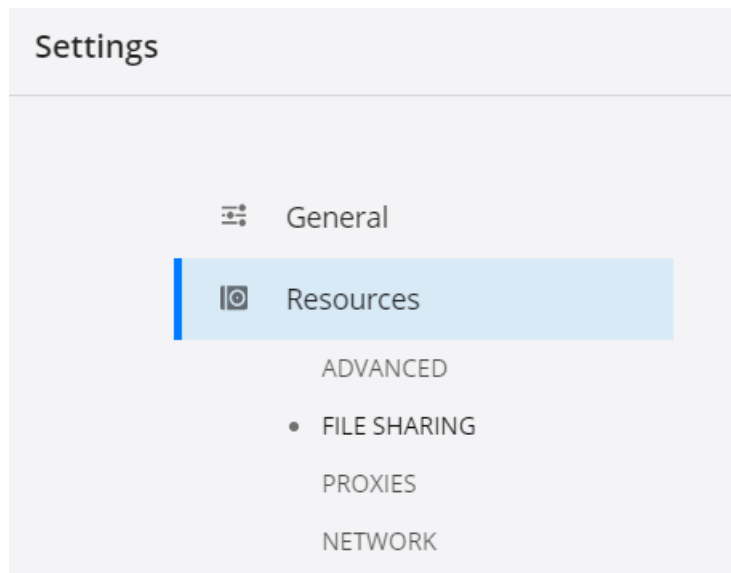
Open Docker for Desktop



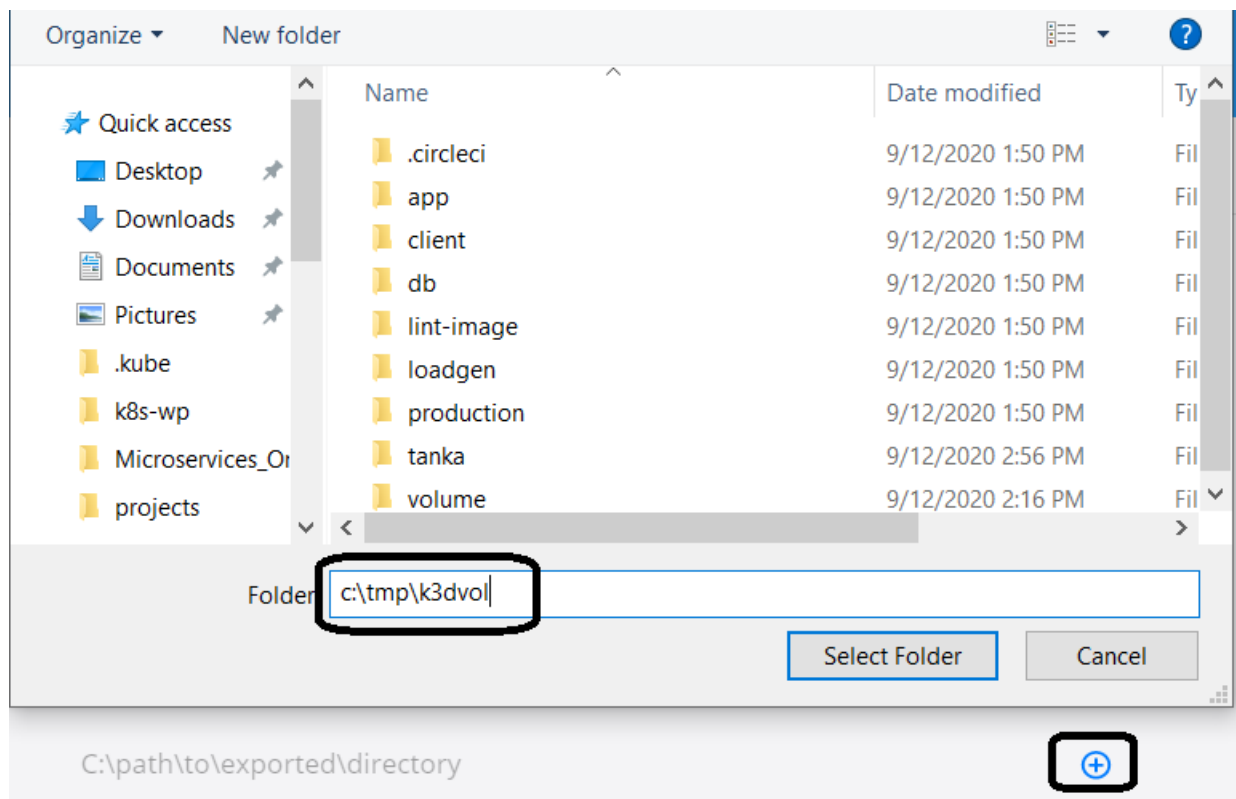
Select the settings icon



Under settings we'll expand the **Resources** and select **FILE SHARING**. From this Settings pane we can select folders that we want to expose as mount points within our Docker containers.

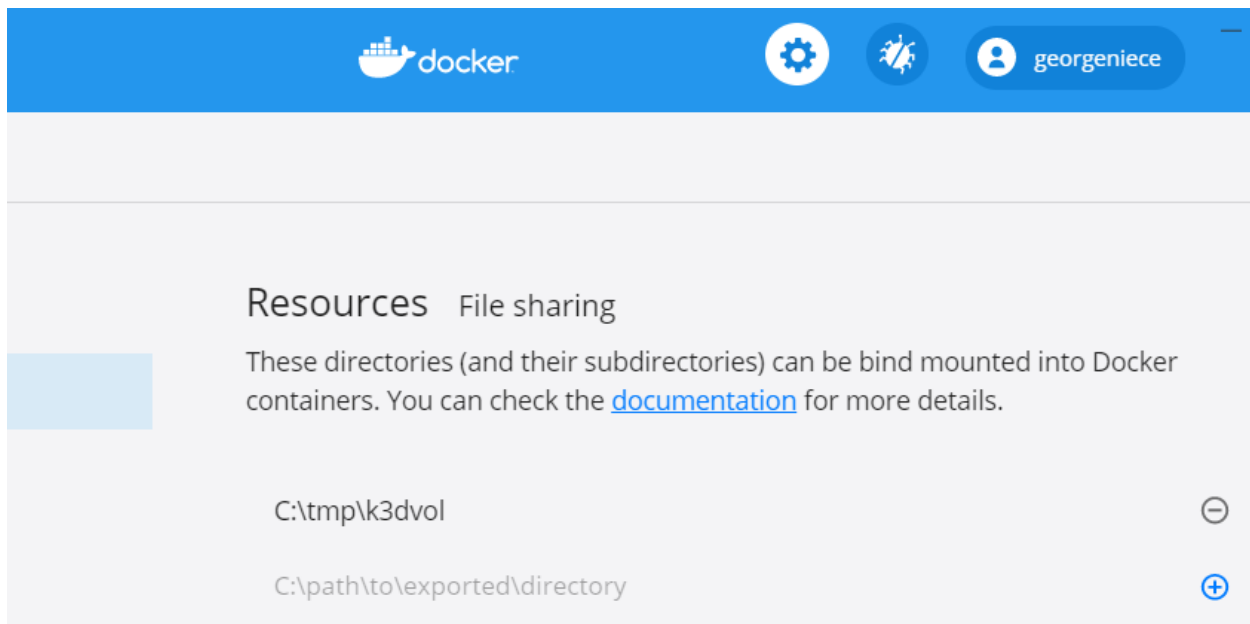


Select the + next to `c:\path\to\exported\directory`



Enter the folder that we created to be bind mounted into our Docker containers. For this experiment that should be `c:\tmp\k3dvol`, then select **Select Folder**

Close, or minimize, the Docker for Desktop dashboard, noting that our new folder is now available in Docker to be mounted for this experiment.



Create the cluster for this experiment, using the volume option to mount our local folder into each node in the cluster. For this cluster we'll expose port 80 against the load balancer and set our agents to 1. Servers option defaults to 1 when not specified.

For Windows:

```
$> k3d cluster create "k3d-cluster" --volume /c/tmp/k3dvol:/tmp/k3dvol --port "80:80@loadbalancer" --agents 2
```

For MacOS:

```
$> mkdir /tmp/k3dvol
```

```
$> k3d cluster create "k3d-cluster" --volume /tmp/k3dvol:/tmp/k3dvol --port "80:80@loadbalancer" --agents 2
```

Note: If you receive an error ensure that you've created the folder for the Persistent Volume (PV) that we're using a Persistent Volume Claim (PVC) for or you'll see an error like the following.

```
[33mWARN[0m[0000] Failed to stat file/directory/named volume that you're trying to mount: '/tmp/k3dvol' in '/tmp/k3dvol:/tmp/k3dvol' -> Please make sure it exists
```

If the volume is connected correctly you should see something similar to the following

```
[33mWARN[0m[0000] No node filter specified
[36mINFO[0m[0000] Created network 'k3d-k3d-cluster'
[36mINFO[0m[0000] Created volume 'k3d-k3d-cluster-images'
[36mINFO[0m[0001] Creating node 'k3d-k3d-cluster-server-0'
[36mINFO[0m[0001] Creating node 'k3d-k3d-cluster-agent-0'
```

```
[36mINFO[0m[0001] Creating node 'k3d-k3d-cluster-agent-1'
[36mINFO[0m[0002] Creating LoadBalancer 'k3d-k3d-cluster-serverlb'
[36mINFO[0m[0008] Cluster 'k3d-cluster' created successfully!
[36mINFO[0m[0008] You can now use it like this:
kubectl cluster-info
```

Set our **KUBECONFIG_FILE** environment variable to the file we'll load our k8s configuration for kubectl usages

```
set KUBECONFIG_FILE=.\kube\k3d-cluster
```

Put our cluster configuration for k3d-cluster into our file

```
k3d kubeconfig get k3d-cluster > %KUBECONFIG_FILE%
```

```
type %KUBECONFIG_FILE%
```

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURStLS0tCk1JSUJWekNCL3FBREFnRUNBZ0VTUUFvR
0NDcUdTTTQ5QkFNQ01DTXhJVEFmQmdOVk1JBTU1HR3N6Y3kxelpYSjIKWlhJdFkyRkFNVF
U1T1RZM01qUTRPEFIRncweU1EQTVNRGt4TnpJNE1EbGFGdzB6TURBNU1EY3hOekk0TU
RsYQpNQ014SVRBZk1JN0ZhdGZFRhOFRRHQxTS2pJekFoTUE0R0ExVWREd0VCL3dRRQpBd0IDc
RPVEJaTUJNR0J5cUdTTTQ5QkFNQ01DTXhJVEFmQmdOVk1JBTU1HR3N6Y3kxelpYSjIKWlhJdFkyRkFNVF
U1T1RZM01qUTRPEFIRncweU1EQTVNRGt4TnpJNE1EbGFGdzB6TURBNU1EY3hOekk0TU
RsYQpNQ014SVRBZk1JN0ZhdGZFRhOFRRHQxTS2pJekFoTUE0R0ExVWREd0VCL3dRRQpBd0IDc
ERBUEJnTIZiUk1CQWY4RUJUQUJBUUg1TUFvR0NDcUdTTTQ5QkFNQ01DTXhJVEFmQmdOVk1JBTU1HR3N6Y3kxelpYSjIKWlhJdFkyRkFNVF
U1T1RZM01qUTRPEFIRncweU1EQTVNRGt4TnpJNE1EbGFGdzB6TURBNU1EY3hOekk0TU
RsYQpNQ014SVRBZk1JN0ZhdGZFRhOFRRHQxTS2pJekFoTUE0R0ExVWREd0VCL3dRRQpBd0IDc
JRM0R2JPCkV4OUxudU45eStrTU44M1AKU1pPWWRGMElYNUV2dXgwPQotLS0tLUVORCB
DRVJUSUZJQ0FURStLS0tCg==
    server: https://0.0.0.0:6550
    name: k3d-k3d-cluster
contexts:
- context:
    cluster: k3d-k3d-cluster
    user: admin@k3d-k3d-cluster
    name: k3d-k3d-cluster
current-context: k3d-k3d-cluster
kind: Config
preferences: {}
users:
- name: admin@k3d-k3d-cluster
  user:
    password: dd79f910ebe64a30855bcd38b7425b98
    username: admin
```

set KUBECONFIG=%KUBECONFIG_FILE%

List our clusters to view the

k3d cluster list k3d-cluster

NAME	SERVICES	AGENTS	LOADBALANCER
k3d-cluster	1/1	2/2	true

kubectl cluster-info

Kubernetes master is running at https://0.0.0.0:51472

CoreDNS is running at https://0.0.0.0:51472/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

Metrics-server is running at https://0.0.0.0:51472/api/v1/namespaces/kube-system/services/https:metrics-server:/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

kubectl cluster-info

Unable to connect to the server: dial tcp [::1]:8080: connectex: No connection could be made because the target machine actively refused it.

Troubleshooting Note: If you have an error similar to above when executing **kubectl**, ensure you correctly set the **KUBECONFIG** in previous steps in this experiment.

Review the enhanced listing for the cluster-info

kubectl cluster-info dump

View the information for the exposed traefik loadbalancer IP/hostname

kubectl get svc traefik --namespace kube-system

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
traefik	LoadBalancer	10.43.245.42	172.18.0.2	80:32162/TCP,443:31433/TCP	42m

kubectl describe svc traefik --namespace kube-system | grep Ingress

LoadBalancer Ingress: 172.18.0.2

Open an editor and paste the following yaml file that will create a **busybox** with a simple ping, as well as exposing our local **c:\tmp\k3dvol** as a mount on **/data** within our container. Alternatively, you can wget the file from the labs folder as

\$ **wget** <https://raw.githubusercontent.com/juel-s/microservices-k8s/master/labs/app.yaml>

View or edit the file to paste the contents and review the yaml file we're using for this experiment

notepad app.yaml

or

vim app.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/tmp/k3dvol"
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: echo
spec:
  selector:
    matchLabels:
      app: echo
  strategy:
```

```

type: Recreate
template:
  metadata:
    labels:
      app: echo
  spec:
    volumes:
      - name: task-pv-storage
        persistentVolumeClaim:
          claimName: task-pv-claim
    containers:
      - image: busybox
        name: echo
        volumeMounts:
          - mountPath: "/data"
            name: task-pv-storage
        command: ["ping", "127.0.0.1"]

```

kubectl apply -f app.yaml

```

persistentvolume/task-pv-volume created
persistentvolumeclaim/task-pv-claim created
deployment.apps/echo created

```

kubectl get pv

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
task-pv-volume	1Gi	RWO	Retain	Bound	default/task-pv-claim manual 38s

View our Persistent Volume Claim

kubectl get pvc

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS
AGE					
task-pv-claim	Bound	task-pv-volume	1Gi	RWO	manual 73s

View our active pod

kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
echo-859c44dcc6-pfc7m	1/1	Running	0	118s

Exec a shell into the container

kubectl exec -it echo-859c44dcc6-pfc7m -- sh

Note: In the deprecated syntax we could have left off the double hyphen before the sh command, but with the pace of change in Kubernetes, and tools like k3d/kind, better to try to stay as current as possible

```
$ echo $(hostname)
```

```
$ echo $(hostname) > /data/hostname.txt
```

```
$ cat /data/hostname.txt
```

```
$ exit
```

```
kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP
OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME				
k3d-k3d-cluster-agent-0	Ready	<none>	53m	v1.18.6+k3s1	172.18.0.3	<none>
Unknown	4.19.76-linuxkit	containerd://1.3.3-k3s2				
k3d-k3d-cluster-agent-1	Ready	<none>	53m	v1.18.6+k3s1	172.18.0.4	<none>
Unknown	4.19.76-linuxkit	containerd://1.3.3-k3s2				
k3d-k3d-cluster-server-0	Ready	master	53m	v1.18.6+k3s1	172.18.0.2	<none>
Unknown	4.19.76-linuxkit	containerd://1.3.3-k3s2				

Delete our identified pod, that we'd exec'd into and created our **hostname.txt** file

```
kubectl delete pod/echo-859c44dcc6-pfc7m
```

```
pod "echo-859c44dcc6-pfc7m" deleted
```

For macOS:

```
cat /tmp/k3dvol/hostname.txt
```

For Windows

```
type c:\tmp\k3dvol\hostname.txt
```

```
echo-859c44dcc6-pfc7m
```

This is Kubernetes, and we're using a configuration requiring that pod for our application, so not surprisingly, we'll do another get pods with kubectl and see a new pod created

```
kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE	READINESS GATES					
echo-859c44dcc6-7mnnr	1/1	Running	0	98m	10.42.2.4	k3d-k3d-cluster-server-0
<none>	<none>					

Exec a sh into our new pod, but this time we'll leave off the double hyphen to see the deprecation warning.

```
kubectl exec -it echo-859c44dcc6-7mnnr sh
```

```
/ # cat /data/hostname.txt
```

```
echo-859c44dcc6-pfc7m
```

```
/ # echo $(hostname)
```

```
echo-859c44dcc6-7mnnr
```

```
/ # exit
```

Delete our cluster for this experiment

```
k3d cluster delete k3d-cluster
```

```
[36mINFO[0m[0000] Deleting cluster 'k3d-cluster'
[36mINFO[0m[0000] Deleted k3d-k3d-cluster-serverlb
[36mINFO[0m[0000] Deleted k3d-k3d-cluster-agent-1
[36mINFO[0m[0000] Deleted k3d-k3d-cluster-agent-0
[36mINFO[0m[0000] Deleted k3d-k3d-cluster-server-0
[36mINFO[0m[0000] Deleting cluster network
'f7f0376fbd55c7f4709ad960ad86c6501ed0a05a19a6d9757914370875a76600'
[36mINFO[0m[0001] Deleting image volume 'k3d-k3d-cluster-images'
[36mINFO[0m[0001] Removing cluster details from default kubeconfig...
[36mINFO[0m[0001] Removing standalone kubeconfig file (if there is one)...
[36mINFO[0m[0001] Successfully deleted cluster k3d-cluster!
```

```
k3d cluster list
```

NAME	SERVICES	AGENTS	LOADBALANCER
------	----------	--------	--------------